

EECS 545: Machine Learning

Lecture 3. Linear Regression (part 2)

Honglak Lee


1/22/2024



Outline

- Linear regression review
- Regularized linear regression
- Review on probability
- Maximum likelihood interpretation of linear regression
- Locally-weighted linear regression

Regression, sum of square error



Linear

$$h(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j \phi_j(\mathbf{x})$$

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left\{ h(\mathbf{x}^{(n)}, \mathbf{w}) - y^{(n)} \right\}^2$$

We want to find \mathbf{w} that minimizes $E(\mathbf{w})$ over the training data.

Least squares problem

- Objective function

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left(\sum_{j=0}^{M-1} w_j \phi_j(\mathbf{x}^{(n)}) - y^{(n)} \right)^2$$

M: dimension of feature, N: number of data

$\phi_j(\mathbf{x}^{(n)})$: j-th feature of data

- Two ways to find \mathbf{w} that minimizes $E(\mathbf{w})$

1. Gradient descent

2. Closed-form solution

Least squares problem

- Objective function

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left(\sum_{j=0}^{M-1} w_j \phi_j(\mathbf{x}^{(n)}) - y^{(n)} \right)^2$$

1. Gradient Descent

$$\frac{\partial E(w)}{\partial w_k} = \sum_{n=1}^N \left(\sum_{j=0}^{M-1} w_j \phi_j(\mathbf{x}^{(n)}) - y^{(n)} \right) \phi_k(\mathbf{x}^{(n)})$$

$$\nabla_{\mathbf{w}} E(\mathbf{w}) = \sum_{n=1}^N \left(\mathbf{w}^\top \phi(\mathbf{x}^{(n)}) - y^{(n)} \right) \phi(\mathbf{x}^{(n)})$$

$$\mathbf{w} := \mathbf{w} - \eta \nabla_{\mathbf{w}} E(\mathbf{w})$$

2. Closed-form solution

$$\begin{aligned} E(\mathbf{w}) &= \frac{1}{2} \sum_{n=1}^N \left(\sum_{j=0}^{M-1} w_j \phi_j(\mathbf{x}^{(n)}) - y^{(n)} \right)^2 \\ &= \frac{1}{2} \mathbf{w}^\top \Phi^\top \Phi \mathbf{w} - \mathbf{w}^\top \Phi^\top \mathbf{y} + \frac{1}{2} \mathbf{y}^\top \mathbf{y} \end{aligned}$$

Recap \mathbf{w} : M by 1 $\mathbf{w} = [w_0, w_1, \dots, w_{M-1}]^\top$

\mathbf{y} : N by 1 $\mathbf{y} = [y^{(1)}, y^{(2)}, \dots, y^{(N)}]^\top$

Φ : N by M


$$\Phi = \begin{pmatrix} \phi_0(\mathbf{x}^{(1)}) & \phi_1(\mathbf{x}^{(1)}) & \dots & \phi_{M-1}(\mathbf{x}^{(1)}) \\ \phi_0(\mathbf{x}^{(2)}) & \phi_1(\mathbf{x}^{(2)}) & \dots & \phi_{M-1}(\mathbf{x}^{(2)}) \\ \vdots & \vdots & \ddots & \vdots \\ \phi_0(\mathbf{x}^{(N)}) & \phi_1(\mathbf{x}^{(N)}) & \dots & \phi_{M-1}(\mathbf{x}^{(N)}) \end{pmatrix}$$

Useful trick: Matrix Calculus

- Compute gradient and set gradient to 0
(condition for optimal solution)

Note: Least squared is a convex problem.

$$\nabla_{\mathbf{w}} E(\mathbf{w}) = 0$$



- Need to compute the first derivative in matrix form

Gradient via matrix calculus

- Compute gradient and set to zero

$$\begin{aligned}\nabla_{\mathbf{w}} E(\mathbf{w}) &= \nabla_{\mathbf{w}} \left(\frac{1}{2} \mathbf{w}^\top \Phi^\top \Phi \mathbf{w} - \mathbf{w}^\top \Phi^\top \mathbf{y} + \frac{1}{2} \mathbf{y}^\top \mathbf{y} \right) \\ &= \Phi^\top \Phi \mathbf{w} - \Phi^\top \mathbf{y} \\ &= \mathbf{0}\end{aligned}$$


- Solve the resulting equation (normal equation)

$$\begin{aligned}\Phi^\top \Phi \mathbf{w} &= \Phi^\top \mathbf{y} \\ \mathbf{w}_{\text{ML}} &= (\Phi^\top \Phi)^{-1} \Phi^\top \mathbf{y}\end{aligned}$$

This is the *Moore-Penrose pseudo-inverse*: $\Phi^\dagger = (\Phi^\top \Phi)^{-1} \Phi^\top$
applied to: $\Phi \mathbf{w} \approx \mathbf{y}$


Back to curve-fitting examples

Polynomial Curve Fitting




$$h(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_{M-1}x^{M-1} = \sum_{j=0}^{M-1} w_j x^j$$


0th Order Polynomial




1st Order Polynomial




3rd Order Polynomial



9th Order Polynomial



Over-fitting




Root-Mean-Square (RMS) Error: $E_{\text{RMS}} = \sqrt{2E(\mathbf{w}^*)/N}$

Q: How do we resolve the over-fitting problem?


Data Set Size: $N = 15$

9th Order Polynomial



Data Set Size: $N = 100$

9th Order Polynomial



Increasing data set size can help

Q. How do we choose the degree of polynomial?

Rule of thumb

- If you have a small number of data, then use low order polynomial (small number of features).
 - Otherwise, your model will overfit
- As you obtain more data, you can gradually increase the order of the polynomial (more features).
 - However, your model is still limited by the finite amount of the data available (i.e., the optimal model for finite data cannot be infinite dimensional polynomial).
- Controlling model complexity: **regularization**

Regularized Linear Regression

Back to Polynomial Coefficients

	$M = 0$	$M = 1$	$M = 3$	$M = 9$
w_0^*	0.19	0.82	0.31	0.35
w_1^*		-1.27	7.99	232.37
w_2^*			-25.43	-5321.83
w_3^*			17.37	48568.31
w_4^*			Good	-231639.30
w_5^*				640042.26
w_6^*				-1061800.52
w_7^*				1042400.18
w_8^*				-557682.99
w_9^*				125201.43

Overfitting;

Coefficients are large!

$$h(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_{M-1}x^{M-1}$$

Regularized Least Squares (1)

- Consider the error function:

$$E_D(\mathbf{w}) + \lambda E_W(\mathbf{w})$$

Data term + Regularization term

λ is called the regularization coefficient.

- With the sum-of-squares error function and a quadratic regularizer, we get Penalize large coefficient values


$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left(\mathbf{w}^\top \phi(\mathbf{x}^{(n)}) - y^{(n)} \right)^2 + \frac{\lambda}{2} \|\mathbf{w}\|_2^2$$

New objective function

Definition (L2): $\|\mathbf{w}\|_2^2 = \sum_{j=0}^{M-1} w_j^2$


- Effect of λ

L2 Regularization: $\ln \lambda = 0$




$$M = 9 \quad \tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left(\mathbf{w}^\top \phi(\mathbf{x}^{(n)}) - y^{(n)} \right)^2 + \frac{\lambda}{2} \|\mathbf{w}\|_2^2$$

L2 Regularization: $\ln \lambda = -18$




$$M = 9 \quad \tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left(\mathbf{w}^\top \phi(\mathbf{x}^{(n)}) - y^{(n)} \right)^2 + \frac{\lambda}{2} \|\mathbf{w}\|_2^2$$

“No” L2 Regularization: $\lambda = 0$
 (or when L2 regularization is too small) (or $\ln \lambda \rightarrow -\infty$)



$$M = 9 \quad \tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left(\mathbf{w}^\top \phi(\mathbf{x}^{(n)}) - y^{(n)} \right)^2 + \frac{\lambda}{2} \|\mathbf{w}\|_2^2$$

L2 Regularization: E_{RMS} vs. $\ln \lambda$



$$E_{\text{RMS}} = \sqrt{2E(\mathbf{w}^*)/N}$$

Larger regularization

NOTE: For simplicity of presentation, we divided the data into training set and test set. However, it's **not** legitimate to find the optimal hyperparameter based on the test set. We will talk about legitimate ways of doing this when we cover model selection and cross-validation.

Polynomial Coefficients

	(i.e., $\lambda = 0$)		
	$\ln \lambda = -\infty$	$\ln \lambda = -18$	$\ln \lambda = 0$
w_0^*	0.35	0.35	0.13
w_1^*	232.37	4.74	-0.05
w_2^*	-5321.83	-0.77	-0.06
w_3^*	48568.31	-31.97	-0.05
w_4^*	-231639.30	-3.89	-0.03
w_5^*	640042.26	55.28	-0.02
w_6^*	-1061800.52	41.32	-0.01
w_7^*	1042400.18	-45.95	-0.00
w_8^*	-557682.99	-91.53	0.00
w_9^*	125201.43	72.68	0.01
	Overfitting; Coefficients are large!	Good	Underfitting

Regularized Least Squares (1)

- Consider the error function:

$$E_D(\mathbf{w}) + \lambda E_W(\mathbf{w})$$

Data term + Regularization term

λ is called the regularization coefficient.

- With the sum-of-squares error function and a quadratic regularizer, we get Penalize large coefficient values

$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left(\mathbf{w}^\top \phi(\mathbf{x}^{(n)}) - y^{(n)} \right)^2 + \frac{\lambda}{2} \|\mathbf{w}\|_2^2$$

- Closed-form solution:

$$\mathbf{w}_{\text{ML}} = (\lambda \mathbf{I} + \Phi^\top \Phi)^{-1} \Phi^\top \mathbf{y}$$

Derivation

Objective function

$$\begin{aligned}\tilde{E}(\mathbf{w}) &= \frac{1}{2} \sum_{n=1}^N \left(\mathbf{w}^\top \phi(\mathbf{x}^{(n)}) - y^{(n)} \right)^2 + \frac{\lambda}{2} \|\mathbf{w}\|_2^2 \\ &= \frac{1}{2} \mathbf{w}^\top \Phi^\top \Phi \mathbf{w} - \mathbf{w}^\top \Phi^\top \mathbf{y} + \frac{1}{2} \mathbf{y}^\top \mathbf{y} + \frac{\lambda}{2} \mathbf{w}^\top \mathbf{w}\end{aligned}$$

Compute the gradient and set it zero:

$$\begin{aligned}\nabla_{\mathbf{w}} \tilde{E}(\mathbf{w}) &= \nabla_{\mathbf{w}} \left[\frac{1}{2} \mathbf{w}^\top \Phi^\top \Phi \mathbf{w} - \mathbf{w}^\top \Phi^\top \mathbf{y} + \frac{1}{2} \mathbf{y}^\top \mathbf{y} + \frac{\lambda}{2} \mathbf{w}^\top \mathbf{w} \right] \\ &= \Phi^\top \Phi \mathbf{w} - \Phi^\top \mathbf{y} + \lambda \mathbf{w} \\ &= (\lambda \mathbf{I} + \Phi^\top \Phi) \mathbf{w} - \Phi^\top \mathbf{y} \\ &= 0\end{aligned}$$

$$\mathbf{w}_{\text{ML}} = (\Phi^\top \Phi)^{-1} \Phi^\top \mathbf{y}$$

v.s. Ordinary Least Square



Therefore, we get: $\mathbf{w}_{\text{ML}} = (\lambda \mathbf{I} + \Phi^\top \Phi)^{-1} \Phi^\top \mathbf{y}$

Regularized Least Squares (2)


- With a more general regularizer, we have

$$\frac{1}{2} \sum_{n=1}^N \left(\mathbf{w}^\top \phi(\mathbf{x}^{(n)}) - y^{(n)} \right)^2 + \frac{\lambda}{2} \sum_{j=1}^M |w_j|^q$$


Note: In this lecture, we focus on $q=2$ (L2 regularization), but other values of $q>0$ can be used.



Lasso
“L1 regularization”



Quadratic
“L2 regularization”




plotting curves of (w_1, w_2) where

$$\sum_{j=1}^M |w_j|^q \text{ is a constant. (M=2)}$$

Regularized Least Squares (3)

- Lasso tends to generate sparser solutions than a quadratic regularizer.



Assuming a simple scenario of isotropic data covariance, the optimal solution to L2/L1 regularization is closest point to the original solution (center of the concentric circles) that touches the boundary of the L2/L1 constraint.

Summary: Regularized Linear Regression

- Simple modification of linear regression
- Regularization controls the tradeoff between “fitting error” and “complexity”
 - Small regularization results in complex models (but with risk of overfitting)
 - Large regularization results in simple models (but with risk of underfitting)
- It is important to find an optimal regularization that balances between the two.

Maximum Likelihood interpretation of least squares regression

Review on probability

Probability: Terminology

- **Experiment:** Procedure that yields an outcome
 - E.g., Tossing a coin three times:
 - Outcome: HHH in one trial, HTH in another trial, etc.
- **Sample space:** Set of all possible outcomes in the experiment, denoted as Ω (or S)
 - E.g., for the above example:
 - $\Omega = \{HHH, HHT, HTH, HTT, THH, THT, THT, TTH, TTT\}$
- **Event:** subset of the sample space Ω (i.e., an event is a set consisting of individual outcomes)
 - Event space: Collection of all events, called \mathcal{F} (aka σ -algebra)
 - E.g., Event that # of heads is an even number.
 - $E = \{HHT, HTH, THH, TTT\}$
- **Probability measure:** function (mapping) from events to probability levels. I.e., $P : \mathcal{F} \rightarrow [0, 1]$ (see next slide)
 - Probability that # of heads is an even number: $4/8 = 1/2$.
- **Probability space:** (Ω, \mathcal{F}, P)

Law of Total Probability

$$P(A) \geq 0, \forall A \in \mathcal{F}$$

$$P(\Omega) \geq 1$$

- Law of total probability

$$P(A) = P(A \cap B) + P(A \cap B^C)$$


$$P(A) = \sum_i P(A \cap B_i) \quad \text{Discrete } B_i$$

$$P(A) = \int P(A \cap B_i) dB_i \quad \text{Continuous } B_i$$

Conditional Probability

For events $A, B \in \mathcal{F}$ with $P(B) > 0$, we may write the **conditional probability** of A given B :

$$P(A|B) = \frac{P(A \cap B)}{P(B)}$$



$$P(A|B_1) = 1$$

$$P(A|B_2) = 0.12 / (0.12 + 0.04) = 0.75$$

$$P(A|B_3) = 0 \text{ (disjoint)}$$

$$P(A) = 0.30 + 0.10 + 0.12 = 0.52$$

(the unconditional probability)

Bayes' Rule

Using the chain rule we may see:

$$P(A|B)P(B) = P(A \cap B) = P(B|A)P(A)$$

Rearranging this yields **Bayes' rule**:

$$P(B|A) = \frac{P(A|B)P(B)}{P(A)}$$

Often this is written as:

$$P(B_i|A) = \frac{P(A|B_i)P(B_i)}{\sum_i P(A|B_i)P(B_i)}$$

Where B_i are a partition of Ω (note the bottom is just the law of total probability).

Likelihood Functions

Why is Bayes' so useful in learning? Allows us to compute the posterior of \mathbf{w} given data D :

$$\text{Posterior} \quad p(\mathbf{w}|D) = \frac{p(D|\mathbf{w})p(\mathbf{w})}{p(D)} \quad \text{Prior}$$

Bayes' rule in words:

$$\text{posterior} \propto \text{likelihood} \times \text{prior}$$
$$p(\mathbf{w}|D) \propto p(D|\mathbf{w})p(\mathbf{w})$$


The likelihood function, $p(\mathbf{w} | D)$, is evaluated for observed data D as a function of \mathbf{w} . It expresses how parameter settings \mathbf{w} .

Maximum Likelihood Estimation (MLE)

- Maximum likelihood:
 - choose parameter setting \mathbf{w} that maximizes likelihood function $p(D | \mathbf{w})$.
 - choose the value of \mathbf{w} that maximizes the probability of observed data.
- Cf. MAP (Maximum a posteriori) estimation
 - Equivalent to maximizing $p(\mathbf{w}|D) \propto p(D|\mathbf{w})p(\mathbf{w})$
 - Can compute this using Bayes rule!
 - This will be covered in later lectures

The Gaussian Distribution

$$\mathcal{N}(x|\mu, \sigma^2) = \frac{1}{(2\pi\sigma^2)^{1/2}} \exp \left\{ -\frac{1}{2\sigma^2} (x - \mu)^2 \right\}$$



$$\mathcal{N}(x|\mu, \sigma^2) > 0$$

$$\int_{-\infty}^{\infty} \mathcal{N}(x|\mu, \sigma^2) dx = 1$$

Maximum Likelihood interpretation of least squares regression

MLE for Linear Regression

- Assume a stochastic model:


$$y^{(n)} = \mathbf{w}^\top \phi(\mathbf{x}^{(n)}) + \epsilon \quad \text{where } \epsilon \sim \mathcal{N}(0, \beta^{-1})$$

- This gives a likelihood function:

$$p(y^{(n)} \mid \phi(\mathbf{x}^{(n)}), \mathbf{w}, \beta) = \mathcal{N}(y^{(n)} \mid \mathbf{w}^\top \phi(\mathbf{x}^{(n)}), \beta^{-1})$$

- With input matrix Φ and output matrix \mathbf{y} , the data likelihood is:

$$p(\mathbf{y} \mid \Phi, \mathbf{w}, \beta) = \prod_{n=1}^N \mathcal{N}(y^{(n)} \mid \mathbf{w}^\top \phi(\mathbf{x}^{(n)}), \beta^{-1})$$



Log-likelihood

- Given data likelihood (prev. slide)

$$p(\mathbf{y} \mid \Phi, \mathbf{w}, \beta) = \prod_{n=1}^N \mathcal{N}(y^{(n)} \mid \mathbf{w}^\top \phi(\mathbf{x}^{(n)}), \beta^{-1})$$

- Log likelihood:

$$\log p(\mathbf{y} \mid \Phi, \mathbf{w}, \beta) = \frac{N}{2} \log \beta - \frac{N}{2} \log 2\pi - \beta E_D(\mathbf{w})$$

$$\text{where } E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left(\mathbf{w}^\top \phi(\mathbf{x}^{(n)}) - y^{(n)} \right)^2$$

- Derivation?

Derivation of log-likelihood of p

From $p(y^{(n)} \mid \phi(\mathbf{x}^{(n)}), \mathbf{w}, \beta) = \mathcal{N}(y^{(n)} \mid \mathbf{w}^\top \phi(\mathbf{x}^{(n)}), \beta^{-1})$

$$= \sqrt{\frac{\beta}{2\pi}} \exp \left(-\frac{\beta}{2} \left\| y^{(n)} - \mathbf{w}^\top \phi(\mathbf{x}^{(n)}) \right\|^2 \right)$$


Derive:

$$\begin{aligned} & \log p(y^{(1)}, y^{(2)}, \dots, y^{(N)} \mid \Phi, \mathbf{w}, \beta) \\ &= \log \prod_{n=1}^N \mathcal{N}(y^{(n)} \mid \mathbf{w}^\top \phi(\mathbf{x}^{(n)}), \beta^{-1}) \\ &= \sum_{n=1}^N \log \left(\sqrt{\frac{\beta}{2\pi}} \exp \left(-\frac{\beta}{2} \left\| y^{(n)} - \mathbf{w}^\top \phi(\mathbf{x}^{(n)}) \right\|^2 \right) \right) \\ &= \sum_{n=1}^N \left(\frac{1}{2} \log \beta - \frac{1}{2} \log 2\pi - \frac{\beta}{2} \left\| y^{(n)} - \mathbf{w}^\top \phi(\mathbf{x}^{(n)}) \right\|^2 \right) \\ &= \frac{N}{2} \log \beta - \frac{N}{2} \log 2\pi - \sum_{n=1}^N \frac{\beta}{2} \left\| y^{(n)} - \mathbf{w}^\top \phi(\mathbf{x}^{(n)}) \right\|^2 \end{aligned}$$

Maximum likelihood estimation (MLE)

- Let's maximize the log-likelihood!
- Set the gradient of log-likelihood = 0 (Why?)

$$\nabla_{\mathbf{w}} \log p(y|\Phi, \mathbf{w}, \beta) = \nabla_{\mathbf{w}} \left(\underbrace{\frac{N}{2} \log \beta - \frac{N}{2} \log 2\pi}_{\text{Constant}} - \sum_{n=1}^N \frac{\beta}{2} \|y^{(n)} - \mathbf{w}^\top \phi(\mathbf{x}^{(n)})\|^2 \right)$$



$$\begin{aligned} &= \beta \sum_{n=1}^N \left(y^{(n)} - \underbrace{\mathbf{w}^\top \phi(\mathbf{x}^{(n)})}_{\text{Scalar}} \phi(\mathbf{x}^{(n)}) \right) \\ &= \beta \left(\sum_{n=1}^N y^{(n)} \phi(\mathbf{x}^{(n)}) - \phi(\mathbf{x}^{(n)}) \phi(\mathbf{x}^{(n)})^\top \mathbf{w} \right) = 0 \end{aligned}$$

- In matrix form, $\beta(\Phi^\top \mathbf{y} - \Phi^\top \Phi \mathbf{w}) = 0$


$$\mathbf{w}_{\text{ML}} = (\Phi^\top \Phi)^{-1} \Phi^\top \mathbf{y}$$

- MLE solution is equivalent to OLS solution!

Locally-weighted Linear Regression


Locally weighted linear regression

- Main idea: When predicting $f(\hat{\mathbf{x}})$, give high weights for “neighbors” of $\hat{\mathbf{x}}$.




In locally weighted regression, points are weighted by proximity to the current $\hat{\mathbf{x}}$ in question using a kernel. A regression is then computed using the weighted points.

Regular linear regression vs. locally weighted linear regression



Regular linear regression

$$\sum_{n=1}^N \left(\mathbf{w}^\top \phi(\mathbf{x}^{(n)}) - y^{(n)} \right)^2$$



Locally weighted linear regression

$$\sum_{n=1}^N r^{(n)}(\hat{\mathbf{x}}) \left(\mathbf{w}^\top \phi(\mathbf{x}^{(n)}) - y^{(n)} \right)^2$$

Linear regression vs. Locally-weighted Linear Regression

- A query point $\hat{\mathbf{x}}$, training set $\left\{ \left(\mathbf{x}^{(n)}, y^{(n)} \right) \right\}_{n=1}^N$

- Linear regression

1. Fit \mathbf{w} to minimize $\sum_{n=1}^N \left(\mathbf{w}^\top \phi(\mathbf{x}^{(n)}) - y^{(n)} \right)^2$
2. Predict $\mathbf{w}^\top \phi(\hat{\mathbf{x}})$

- Locally-weighted linear regression

1. Fit \mathbf{w} to minimize $\sum_{n=1}^N r^{(n)}(\hat{\mathbf{x}}) \left(\mathbf{w}^\top \phi(\mathbf{x}^{(n)}) - y^{(n)} \right)^2$
2. Predict $\mathbf{w}^\top \phi(\hat{\mathbf{x}})$

weights are dependent on the query $\hat{\mathbf{x}}$
(i.e., need to solve the optimization for each query value)

Linear regression vs. Locally-weighted Linear Regression

- Locally-weighted linear regression

1. Fit \mathbf{w} to minimize $\sum_{n=1}^N r^{(n)}(\hat{\mathbf{x}}) \left(\mathbf{w}^\top \phi(\mathbf{x}^{(n)}) - y^{(n)} \right)^2$
2. Predict $\mathbf{w}^\top \phi(\hat{\mathbf{x}})$


- Remarks:

Gaussian kernel with kernel width τ

1. Standard choice: $r^{(n)}(\hat{\mathbf{x}}) = \exp \left(-\frac{\|\phi(\mathbf{x}^{(n)}) - \phi(\hat{\mathbf{x}})\|^2}{2\tau^2} \right)$
2. Note that $r^{(n)}(\hat{\mathbf{x}})$ depends on $\hat{\mathbf{x}}$ (query point), and you solve linear regression for each query point $\hat{\mathbf{x}}$
3. The problem can be formulated as a modified version of least squares problem (HW#1)

Locally weighted linear regression

- Choice of kernel width τ matters
 - Requires hyper-parameter tuning



The estimator is minimized when kernel includes as many training points as can be accommodated by the model. Too large a kernel includes points that degrade the fit; too small a kernel neglects points that increase confidence in the fit.

Summary

- L_2 Regularized linear regression $\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left(\mathbf{w}^\top \phi(\mathbf{x}^{(n)}) - y^{(n)} \right)^2 + \frac{\lambda}{2} \|\mathbf{w}\|_2^2$
 - Adding L_2 regularizer
 - Can be solved via closed form (simple modification of the original linear regression)
 - penalizes complex solutions (with high weights)
- Maximum likelihood interpretation of linear regression
 - Linear regression can be interpreted as performing MLE assuming the Gaussian noise distribution for targets
- Locally-weighted linear regression

Any feedback (about lecture, slide, homework, project, etc.)?

(via **anonymous** google form: <https://forms.gle/99jeftYTaozJvCEF8>)



Change Log of lecture slides:

https://docs.google.com/document/d/e/2PACX-1vRKx40eOJKACqrKWraio0AmlFS1_xBMINuWcc-jzpfo-ySj_gBugTVdfHy8v4HDmqDJ3b3TvaW1FVuH/pub