

EECS 545: Machine Learning

Lecture 4. Classification

Honglak Lee

1/24/2024




Outline

- Logistic regression
- Newton's method
- K-nearest neighbors (KNN)


Supervised learning: classification

Supervised learning

- Goal:
 - Given data X in feature space and labels Y
 - Learn to predict Y from X
- Labels could be discrete or continuous
 - Discrete-valued labels: classification (today's topic)
 - Continuous-valued labels: regression



classification



regression

Classification problem

- The task of classification:
 - Given an input vector \mathbf{x} , assign it to one of K distinct classes C_k where $k = 1, \dots, K$
- Representing the assignment:
 - For $K = 2$:
 - $y = 1$ means that \mathbf{x} is in C_1
 - $y = 0$ means that \mathbf{x} is in C_2 .
 - (Sometimes, $y = -1$ can be used depending on algorithms)
- For $K > 2$:
 - Use 1-of- K coding
 - e.g., $\mathbf{y} = (0, 1, 0, 0, 0)^T$ means that \mathbf{x} is in C_2 .
 - (This works for $K = 2$ as well)

Classification problem



- Training: train a classifier $h(\mathbf{x})$ from training data
 - Training data $\{ (x^{(1)}, y^{(1)}) , (x^{(2)}, y^{(2)}) , \dots , (x^{(N)}, y^{(N)}) \}$
- Testing (evaluation):
 - testing data: $h \left(x_{\text{test}}^{(1)} \right) , h \left(x_{\text{test}}^{(2)} \right) , \dots , h \left(x_{\text{test}}^{(N')} \right)$
 - The learning algorithm produces predictions
 - 0-1 loss: Classification error $= \frac{1}{N'} \sum_{j=1}^{N'} \mathbb{I} \left[h \left(x_{\text{test}}^{(j)} \right) \neq y_{\text{test}}^{(j)} \right]$

Logistic regression

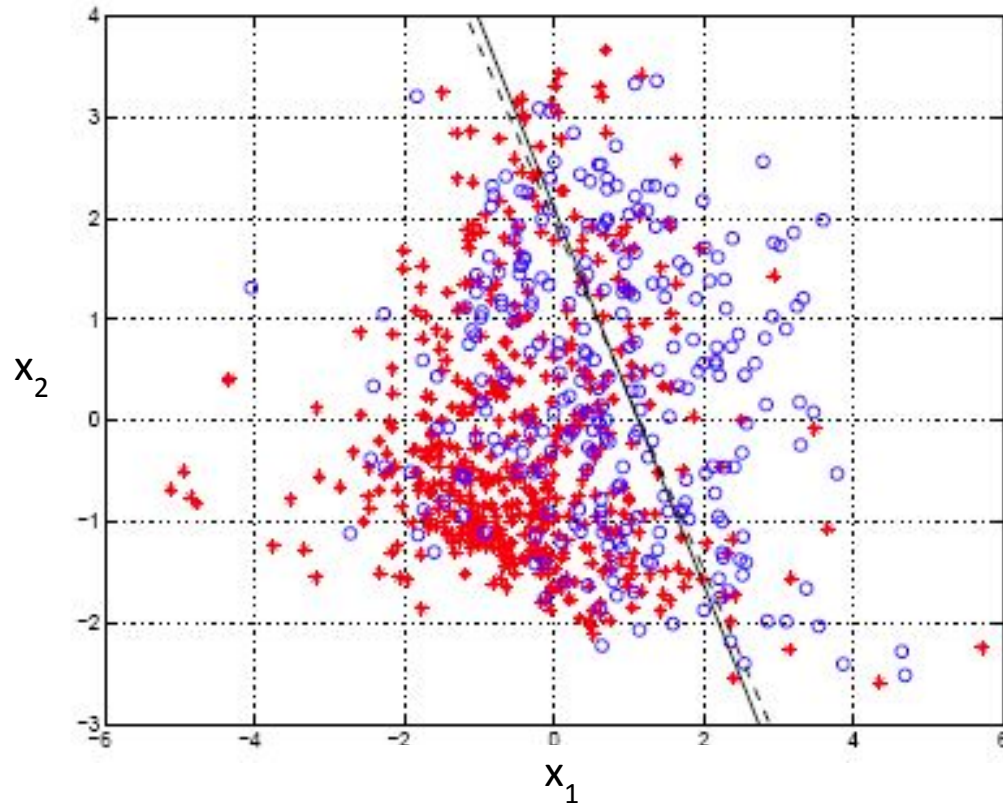
Probabilistic discriminative models

- Model decision boundary as a function of input \mathbf{x}
 - Learn $P(C_k | \mathbf{x})$ over data (e.g., maximum likelihood)
 - Directly predict class labels from inputs
- Next class: we will cover probabilistic generative models
 - Learn $P(C_k, \mathbf{x})$ over data (maximum likelihood) and then use Bayes' rule to predict $P(C_k | \mathbf{x})$

Example (1-dim. case)



Example (2-dim. case)



Logistic regression

- Models the class posterior using a sigmoid applied to a linear function of the feature vector:


$$p(C_1|\phi) = h(\phi) = \sigma(\mathbf{w}^\top \phi(\mathbf{x}))$$

- We can solve the parameter \mathbf{w} by maximizing the likelihood of the training data

Sigmoid and logit functions

- The *logistic sigmoid* function is:

$$\sigma(a) = \frac{1}{1 + \exp(-a)} = \frac{\exp(a)}{1 + \exp(a)}$$



- Its inverse is the *logit* function (aka log odds ratio):

$$a = \ln \left(\frac{\sigma}{1 - \sigma} \right)$$

- Generalizes to *normalized exponential*, or *softmax*

$$p_i = \frac{\exp(q_i)}{\sum_j \exp(q_j)}$$

Class-conditional probability (for a single example)

- Depending on the label y , the conditional probability of y given \mathbf{x} is defined as:

$$P(y = 1|\mathbf{x}, \mathbf{w}) = \sigma(\mathbf{w}^\top \phi(\mathbf{x}))$$

$$P(y = 0|\mathbf{x}, \mathbf{w}) = 1 - \sigma(\mathbf{w}^\top \phi(\mathbf{x}))$$

- Therefore we can write both cases compactly as:

$$P(y|\mathbf{x}, \mathbf{w}) = \sigma(\mathbf{w}^\top \phi(\mathbf{x}))^y (1 - \sigma(\mathbf{w}^\top \phi(\mathbf{x})))^{1-y}$$

Likelihood function (of logistic regression)

- The likelihood of Data $\{(\phi(\mathbf{x}^{(n)}), y^{(n)})\}$, where $y^{(n)} \in \{0, 1\}$

$$P(D|\mathbf{w}) = \prod_{i=1}^N P(\mathbf{x}^{(i)}, y^{(i)}|\mathbf{w}) \quad \text{IID (Independent Identical Distribution)}$$

Definition of
conditional
probability



$$= \prod_{i=1}^N P(y^{(i)}|\mathbf{x}^{(i)}, \mathbf{w}) \underbrace{P(\mathbf{x}^{(i)}|\mathbf{w})}_{=P(\mathbf{x}^{(i)})}$$

P(x) does not depend on w

$$\propto \prod_{i=1}^N P(y^{(i)}|\mathbf{x}^{(i)}, \mathbf{w}) \longrightarrow P(\mathbf{y}|\mathbf{X}, \mathbf{w})$$

Compact notation: Technically speaking, this is
(conditional) likelihood of y given X

Logistic regression

- For a data set $\{(\phi(\mathbf{x}^{(n)}), y^{(n)})\}$, where $y^{(n)} \in \{0, 1\}$ the likelihood function is

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \prod_{n=1}^N (h^{(n)})^{y^{(n)}} (1 - h^{(n)})^{1-y^{(n)}}$$

note: $h(\mathbf{x})$ is the hypothesis function, $\sigma(\mathbf{x})$ is the specific hypothesis for logistic regression

where

$$h^{(n)} = p(C_1|\phi(\mathbf{x}^{(n)})) = \sigma(\mathbf{w}^\top \phi(\mathbf{x}^{(n)}))$$

- Define a loss function $E(\mathbf{w}) = -\log p(\mathbf{y}|\mathbf{X}, \mathbf{w})$
 - Minimizing $E(\mathbf{w})$ maximizes likelihood

Derivation

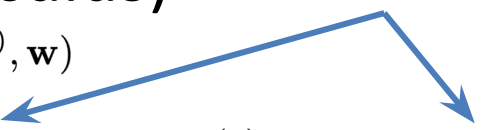
- $\log P(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \sum_{n=1}^N y^{(n)} \log h^{(n)} + (1 - y^{(n)}) \log(1 - h^{(n)})$
- Gradient (matrix calculus)

$$\begin{aligned} & \nabla_{\mathbf{w}} \log P(\mathbf{y}|\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}, \mathbf{w}) \\ &= \sum_{n=1}^N \nabla_{\mathbf{w}} \left(y^{(n)} \log h(\mathbf{x}^{(n)}, \mathbf{w}) + (1 - y^{(n)}) \log(1 - h(\mathbf{x}^{(n)}, \mathbf{w})) \right) \end{aligned}$$

Derivation

- $\log P(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \sum_{n=1}^N y^{(n)} \log h^{(n)} + (1 - y^{(n)}) \log(1 - h^{(n)})$
- **Gradient (matrix calculus)**
$$\nabla_{\mathbf{w}} \log P(\mathbf{y}|\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}, \mathbf{w})$$
$$= \sum_{n=1}^N \nabla_{\mathbf{w}} \left(y^{(n)} \log h(\mathbf{x}^{(n)}, \mathbf{w}) + (1 - y^{(n)}) \log(1 - h(\mathbf{x}^{(n)}, \mathbf{w})) \right)$$


$h(\mathbf{x}^{(n)}, \mathbf{w}) \triangleq \sigma(\mathbf{w}^\top \phi(\mathbf{x}^{(n)})) \triangleq \sigma^{(n)}$



Derivation

- $\log P(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \sum_{n=1}^N y^{(n)} \log h^{(n)} + (1 - y^{(n)}) \log(1 - h^{(n)})$
- Gradient (matrix calculus)** $h(\mathbf{x}^{(n)}, \mathbf{w}) \triangleq \sigma(\mathbf{w}^\top \phi(\mathbf{x}^{(n)})) \triangleq \sigma^{(n)}$

$$\begin{aligned} & \nabla_{\mathbf{w}} \log P(\mathbf{y}|\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}, \mathbf{w}) \\ &= \sum_{n=1}^N \nabla_{\mathbf{w}} \left(y^{(n)} \log h(\mathbf{x}^{(n)}, \mathbf{w}) + (1 - y^{(n)}) \log(1 - h(\mathbf{x}^{(n)}, \mathbf{w})) \right) \\ &= \sum_{n=1}^N \left(y^{(n)} \frac{\sigma^{(n)}(1 - \sigma^{(n)})}{\sigma^{(n)}} - (1 - y^{(n)}) \frac{\sigma^{(n)}(1 - \sigma^{(n)})}{1 - \sigma^{(n)}} \right) \nabla_{\mathbf{w}} (\mathbf{w}^\top \phi(\mathbf{x}^{(n)})) \end{aligned}$$


$$\frac{\partial}{\partial s} \sigma(s) = \frac{\partial}{\partial s} \left(\frac{1}{1 + \exp(-s)} \right) = \sigma(s)(1 - \sigma(s))$$

Derivation

- $\log P(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \sum_{n=1}^N y^{(n)} \log h^{(n)} + (1 - y^{(n)}) \log(1 - h^{(n)})$
- **Gradient (matrix calculus)** $h(\mathbf{x}^{(n)}, \mathbf{w}) \triangleq \sigma(\mathbf{w}^\top \phi(\mathbf{x}^{(n)})) \triangleq \sigma^{(n)}$

$$\begin{aligned} & \nabla_{\mathbf{w}} \log P(\mathbf{y}|\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}, \mathbf{w}) \\ &= \sum_{n=1}^N \nabla_{\mathbf{w}} \left(y^{(n)} \log h(\mathbf{x}^{(n)}, \mathbf{w}) + (1 - y^{(n)}) \log(1 - h(\mathbf{x}^{(n)}, \mathbf{w})) \right) \\ &= \sum_{n=1}^N \left(y^{(n)} \frac{\sigma^{(n)}(1 - \sigma^{(n)})}{\sigma^{(n)}} - (1 - y^{(n)}) \frac{\sigma^{(n)}(1 - \sigma^{(n)})}{1 - \sigma^{(n)}} \right) \nabla_{\mathbf{w}} (\mathbf{w}^\top \phi(\mathbf{x}^{(n)})) \\ &= \sum_{n=1}^N \left(y^{(n)}(1 - \sigma^{(n)}) - (1 - y^{(n)})\sigma^{(n)} \right) \nabla_{\mathbf{w}} (\mathbf{w}^\top \phi(\mathbf{x}^{(n)})) \end{aligned}$$

Derivation

- $\log P(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \sum_{n=1}^N y^{(n)} \log h^{(n)} + (1 - y^{(n)}) \log(1 - h^{(n)})$
- Gradient (matrix calculus)** $h(\mathbf{x}^{(n)}, \mathbf{w}) \triangleq \sigma(\mathbf{w}^\top \phi(\mathbf{x}^{(n)})) \triangleq \sigma^{(n)}$

$$\begin{aligned} & \nabla_{\mathbf{w}} \log P(\mathbf{y}|\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}, \mathbf{w}) \\ &= \sum_{n=1}^N \nabla_{\mathbf{w}} \left(y^{(n)} \log h(\mathbf{x}^{(n)}, \mathbf{w}) + (1 - y^{(n)}) \log(1 - h(\mathbf{x}^{(n)}, \mathbf{w})) \right) \\ &= \sum_{n=1}^N \left(y^{(n)} \frac{\sigma^{(n)}(1 - \sigma^{(n)})}{\sigma^{(n)}} - (1 - y^{(n)}) \frac{\sigma^{(n)}(1 - \sigma^{(n)})}{1 - \sigma^{(n)}} \right) \nabla_{\mathbf{w}} (\mathbf{w}^\top \phi(\mathbf{x}^{(n)})) \\ &= \sum_{n=1}^N \left(y^{(n)}(1 - \sigma^{(n)}) - (1 - y^{(n)})\sigma^{(n)} \right) \nabla_{\mathbf{w}} (\mathbf{w}^\top \phi(\mathbf{x}^{(n)})) \\ &= \sum_{n=1}^N \left(y^{(n)} - \sigma^{(n)} \right) \phi(\mathbf{x}^{(n)}) \end{aligned}$$

Logistic regression: gradient descent

- Taking the gradient of $E(\mathbf{w})$ gives us

$$\nabla E(\mathbf{w}) = \sum_{n=1}^N (h^{(n)} - y^{(n)}) \phi(\mathbf{x}^{(n)})$$

- Recall

$$h^{(n)} = p(C_1 | \phi(\mathbf{x}^{(n)})) = \sigma(\mathbf{w}^\top \phi(\mathbf{x}^{(n)}))$$

- This is essentially the same gradient expression that appeared in linear regression with least-squares.
- Note the error term between model prediction and target value:
 - Logistic regression: $h^{(n)} - y^{(n)} = \sigma(\mathbf{w}^\top \phi(\mathbf{x}^{(n)})) - y^{(n)}$
 - Cf. Linear regression: $h^{(n)} - y^{(n)} = \mathbf{w}^\top \phi(\mathbf{x}^{(n)}) - y^{(n)}$

Newton's method

- Goal: Minimizing a general function $E(\mathbf{w})$ (one-dimensional case)

- Approach: solve for

$$f(\mathbf{w}) = \frac{\partial E(\mathbf{w})}{\partial \mathbf{w}} = 0$$

- So, how to solve this problem?


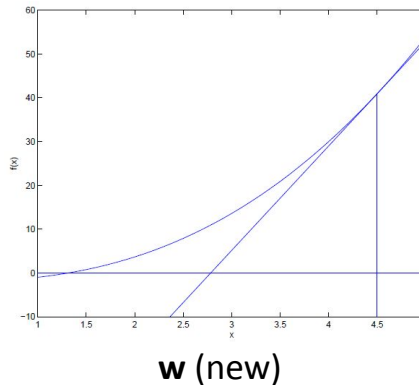
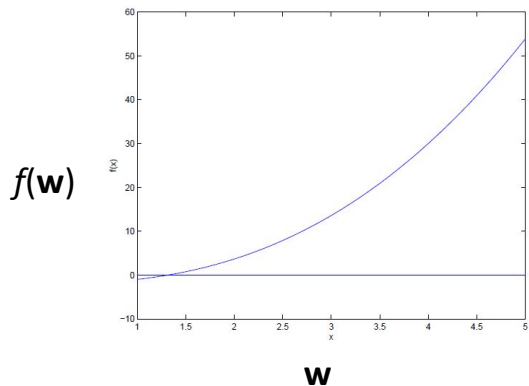
- Newton's method (aka Newton-Raphson method)

- Repeat until convergence:

$$\mathbf{w} := \mathbf{w} - \frac{f(\mathbf{w})}{f'(\mathbf{w})}$$

Newton's method

- Interactively solve until we get $f(w) = 0$.



- Geometric intuition:

$$w := w - \frac{f(w)}{f'(w)}$$

Current value

"Slope"

Newton's method

- Now we want to minimize $E(\mathbf{w})$
 - Convert $E'(\mathbf{w}) = f(\mathbf{w})$
 - Repeat until convergence

$$\mathbf{w} := \mathbf{w} - \frac{E'(\mathbf{w})}{E''(\mathbf{w})}$$

Newton update
when w is a scalar

Newton's method

- Now we want to minimize $E(\mathbf{w})$
 - Convert $E'(\mathbf{w}) = f(\mathbf{w})$
 - Repeat until convergence

$$\mathbf{w} := \mathbf{w} - \frac{E'(\mathbf{w})}{E''(\mathbf{w})}$$

Newton update
when w is a scalar

- This method can be extended to the multivariate case:

$$\mathbf{w} := \mathbf{w} - H^{-1} \nabla_{\mathbf{w}} E$$

Newton update
when w is a vector

where \mathbf{H} is a Hessian matrix evaluated at \mathbf{w}

$$H_{ij}(\mathbf{w}) = \frac{\partial^2 E(\mathbf{w})}{\partial \mathbf{w}_i \partial \mathbf{w}_j}$$

- Note: for linear regression, the Hessian is $\Phi^\top \Phi$

Logistic regression

- Recall: for linear regression, least-squares has a closed-form solution: $\mathbf{w}_{\text{ML}} = (\Phi^\top \Phi)^{-1} \Phi^\top \mathbf{y}$
- This generalizes to weighted-least-squares with an $N \times N$ diagonal weight matrix \mathbf{R} .

$$\mathbf{w}_{\text{WLS}} = (\Phi^\top \mathbf{R} \Phi)^{-1} \Phi^\top \mathbf{R} \mathbf{y}$$

- For logistic regression, however, $h(\mathbf{x}, \mathbf{w})$ is non-linear, and there is no closed-form solution. Must iterate (i.e. repeatedly apply Newton steps).

Iterative solution

- Apply Newton-Raphson method to iterate to a solution \mathbf{w} for $\nabla E(\mathbf{w}) = 0$
- This involves least-squares with weights \mathbf{R} :

$$R_{nn} = h^{(n)}(1 - h^{(n)})$$

- Since \mathbf{R} depends on \mathbf{w} (and vice versa), we get *iterative reweighted least squares* (IRLS)

where $\mathbf{w}^{(\text{new})} = (\Phi^\top \mathbf{R} \Phi)^{-1} \Phi^\top \mathbf{R} \mathbf{z}$

$$\mathbf{z} = \Phi \mathbf{w}^{(\text{old})} - \mathbf{R}^{-1}(\mathbf{h} - \mathbf{y})$$

K-nearest neighbor classification

K-nearest neighbors

- Training method:
 - Save the training examples (no sophisticated learning)
- At prediction (testing) time:
 - Given a test (query) example \mathbf{x} , find the K training examples that are *closest* to \mathbf{x} .



$$\text{KNN}(\mathbf{x}) = \{ (\mathbf{x}^{(1)'}, y^{(1)'}) , (\mathbf{x}^{(2)'}, y^{(2)'}) , ..., (\mathbf{x}^{(K)'}, y^{(K)'}) \}$$

- Predict the most frequent class among all y 's from $\text{KNN}(\mathbf{x})$.


$$h(\mathbf{x}) = \arg \max_y \sum_{(\mathbf{x}', y') \in \text{KNN}(\mathbf{x})} \mathbb{I}[y' = y] \quad \text{“majority vote”}$$

- Note: this function can be applied to regression!

K-nearest neighbors for classification



K-nearest neighbors for classification




- Larger K leads to a smoother decision boundary (bias-variance trade-off)
- Classification performance generally improves as N (training set size) increases
- For $N \Rightarrow \infty$, the error rate of the 1-nearest-neighbor classifier is never more than twice the optimal error (obtained from the true conditional class distributions). See ESL CH 13.3.

Factors (hyperparameters) affecting KNN

- Distance metric $D(\mathbf{x}, \mathbf{x}')$
 - How to define distance between two examples \mathbf{x} and \mathbf{x}' ?
- The value of K
 - K determines how much we “smooth out” the prediction

What is the decision boundary?

Voronoi diagram: Euclidean (L_2) distance




Note: Each region corresponds kNN's prediction when $K=1$



i.e. prediction is the same as the corresponding training sample's label in each region (training sample is visualized as dot).

Dependence on distance metric (L^q norm)



Distance between i-th and j-th example: $\sqrt[q]{\sum_l (x_l^{(i)} - x_l^{(j)})^q}$




knn (K=1): l1 Distance



knn (K=1): l2 Distance



knn (K=1): linf Distance



KNN: classification vs regression

- We can formulate KNN into regression/classification
- For classification, where the label y is categorical, we take the “majority vote” over target labels.

$$h(\mathbf{x}) = \arg \max_y \sum_{(\mathbf{x}', y') \in \text{KNN}(\mathbf{x})} \mathbb{I}[y' = y]$$

- For regression, where the label y is real-valued numbers, we take “average” over target labels.

$$h(\mathbf{x}) = \frac{1}{k} \sum_{(\mathbf{x}', y') \in \text{KNN}(\mathbf{x})} y'$$

Advantage/disadvantages of KNN methods

- Advantage:
 - Very simple and flexible (no assumption on distribution)
 - Effective (e.g. for low dimensional inputs)
- Disadvantages:
 - Expensive: need to remember (store) and search through all the training data for every prediction
 - Curse of dimensionality: in high dimensions, all points are far
 - Not robust to irrelevant features: if \mathbf{x} has irrelevant/noisy features, then distance function does not reflect similarity between examples

Concept check

- How are labels represented in multiclass classification problems?
- What is the motivation for using Newton's method for optimization in logistic regression?
- What does increasing K do for the results from KNN?

Any feedback (about lecture, slide, homework, project, etc.)?

(via anonymous google form: <https://forms.gle/99jeftYTaozJvCEF8>)



Change Log of lecture slides:

https://docs.google.com/document/d/e/2PACX-1vRKx40eOJKACqrKWraio0AmIFS1_xBMINuWcc-jzpfo-ySj_gBuqTVdfHy8v4HDmqDJ3b3TVAW1FVuH/pub